

Reverse Engineering

Week 2: Memory Structures, Dynamic RE, Linux syscalls/ABI



Sergey Bratus, Travis Goodspeed and Ryan Speers -- Dartmouth College -- Winter 2022

Overview

What & Why

- Today will mostly be us stepping through examples followed by a lab to do on your own
- The goal is to get familiar with a few common utilities and to give you a chance to ask any questions
- The homework for today will just be doing what you learn today on a binary that you hadn't seen before

Motivation

Cool RE

- <https://reactos.org>
- Windows is “closed source”
- This group is slowly reverse engineering Windows components and rewriting them in C/C++.

nasm - netwide assembler

What & Why

- Used to compile assembly programs into binary blobs or even ELF files
- I primarily use it to compile snippets of assembly to see the corresponding machine code
- Really useful for something you may learn later called shellcode.
- Connect to a babylon server and we will start debugging together

nasm - netwide assembler

Write the code

```
BITS 64
```

```
_start:
```

```
    mov rax, rdi
```

```
    ret
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

```
~
```

nasm - netwide assembler

Compile it

```
f00543d@babylon1:~/recourse/week2$ nasm -f bin assembly.asm  
f00543d@babylon1:~/recourse/week2$
```

nasm - netwide assembler

Check it out

```
f00543d@babylon1:~/recourse/week2$ nasm -f bin assembly.asm
f00543d@babylon1:~/recourse/week2$ ndisasm -b 64 assembly
00000000 4889F8          mov rax,rdi
00000003 C3             ret
f00543d@babylon1:~/recourse/week2$ █
```

nasm - netwide assembler

Format it for use

```
f00543d@babylon1:~/recourse/week2$ nasm -f bin assembly.asm
f00543d@babylon1:~/recourse/week2$ ndisasm -b 64 assembly
00000000  4889F8          mov rax,rdi
00000003  C3             ret
f00543d@babylon1:~/recourse/week2$ xxd -i assembly
unsigned char assembly[] = {
    0x48, 0x89, 0xf8, 0xc3
};
unsigned int assembly_len = 4;
f00543d@babylon1:~/recourse/week2$
```


nasm - netwide assembler

Compile it

```
f00543d@babylon1:~/recourse/week2$ nasm -f bin assembly.asm
f00543d@babylon1:~/recourse/week2$ ndisasm -b 64 assembly
00000000  4889F8          mov rax,rdi
00000003  C3             ret
f00543d@babylon1:~/recourse/week2$ xxd -i assembly
unsigned char assembly[] = {
    0x48, 0x89, 0xf8, 0xc3
};
unsigned int assembly_len = 4;
f00543d@babylon1:~/recourse/week2$ gcc test_exec.c -z execstack -o test_exec
f00543d@babylon1:~/recourse/week2$
```

nasm - netwide assembler

Run it

```
f00543d@babylon1:~/recourse/week2$ nasm -f bin assembly.asm
f00543d@babylon1:~/recourse/week2$ ndisasm -b 64 assembly
00000000  4889F8          mov rax,rdi
00000003  C3              ret
f00543d@babylon1:~/recourse/week2$ xxd -i assembly
unsigned char assembly[] = {
    0x48, 0x89, 0xf8, 0xc3
};
unsigned int assembly_len = 4;
f00543d@babylon1:~/recourse/week2$ gcc test_exec.c -z execstack -o test_exec
f00543d@babylon1:~/recourse/week2$ ./test_exec 1337
Return value: 1337
f00543d@babylon1:~/recourse/week2$
```

gdb - Gnu Debugger

What & Why

- Can attach to an already running process or launch one itself
- Allows you to see the program state and step through a concrete execution
- Sometimes it really helps to see an actual execution when dealing with complex code
- Connect to a babylon server and we will start debugging together

gdb - Gnu Debugger

Launch

- Copy the `hello_world` binary into your folder and in your terminal run
 - `gdb hello_world`

```
f00543d@babylon1:~/recourse$ gdb ./hello_world
GNU gdb (Ubuntu 8.1.1-0ubuntu1) 8.1.1
Copyright (C) 2018 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.  Type "show copying"
and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.
For help, type "help".
Type "apropos word" to search for commands related to "word"...
Reading symbols from ./hello_world...(no debugging symbols found)...done.
(gdb) █
```


gdb - Gnu Debugger

A look around

- The application isn't currently running but there are a few things that we can do.
- Let's check out the disassembly for main.
 - run 'disassemble main'
 - The output you see is in AT&T Gas syntax
 - To see it as INTEL syntax run
 - 'set disassembly-flavor intel'

```
(gdb) disassemble main
Dump of assembler code for function main:
0x000000000000017d <+0>:      endbr64
0x0000000000000181 <+4>:      push   %rbp
0x0000000000000182 <+5>:      mov    %rsp,%rbp
0x0000000000000185 <+8>:      sub    $0x10,%rsp
0x0000000000000189 <+12>:     mov    %edi,-0x4(%rbp)
0x000000000000018c <+15>:     mov    %rsi,-0x10(%rbp)
0x0000000000000190 <+19>:     cmpl   $0x1,-0x4(%rbp)
0x0000000000000194 <+23>:     jle    0x11ab <main+46>
0x0000000000000196 <+25>:     mov    -0x10(%rbp),%rax
0x000000000000019a <+29>:     add    $0x8,%rax
0x000000000000019e <+33>:     mov    (%rax),%rax
0x00000000000001a1 <+36>:     mov    %rax,%rdi
0x00000000000001a4 <+39>:     callq  0x1149 <print_string>
0x00000000000001a9 <+44>:     jmp    0x11b7 <main+58>
0x00000000000001ab <+46>:     lea    0xe5e(%rip),%rdi      # 0x2010
0x00000000000001b2 <+53>:     callq  0x1149 <print_string>
0x00000000000001b7 <+58>:     mov    $0x0,%eax
0x00000000000001bc <+63>:     leaveq
0x00000000000001bd <+64>:     retq
End of assembler dump.
(gdb) █
```

gdb - Gnu Debugger

A look around

```
(gdb) set disassembly-flavor intel
(gdb) disassemble main
Dump of assembler code for function main:
   0x0000000000000117d <+0>:      endbr64
   0x00000000000001181 <+4>:      push    rbp
   0x00000000000001182 <+5>:      mov     rbp, rsp
   0x00000000000001185 <+8>:      sub     rsp, 0x10
   0x00000000000001189 <+12>:     mov     DWORD PTR [rbp-0x4], edi
   0x0000000000000118c <+15>:     mov     QWORD PTR [rbp-0x10], rsi
   0x00000000000001190 <+19>:     cmp     DWORD PTR [rbp-0x4], 0x1
   0x00000000000001194 <+23>:     jle     0x11ab <main+46>
   0x00000000000001196 <+25>:     mov     rax, QWORD PTR [rbp-0x10]
   0x0000000000000119a <+29>:     add     rax, 0x8
   0x0000000000000119e <+33>:     mov     rax, QWORD PTR [rax]
   0x000000000000011a1 <+36>:     mov     rdi, rax
   0x000000000000011a4 <+39>:     call    0x1149 <print_string>
   0x000000000000011a9 <+44>:     jmp     0x11b7 <main+58>
   0x000000000000011ab <+46>:     lea     rdi, [rip+0xe5e]          # 0x2010
   0x000000000000011b2 <+53>:     call    0x1149 <print_string>
   0x000000000000011b7 <+58>:     mov     eax, 0x0
   0x000000000000011bc <+63>:     leave
   0x000000000000011bd <+64>:     ret
End of assembler dump.
(gdb) █
```

gdb - Gnu Debugger

A look around

- You will see that the address of the first instruction is 0x117d. Remember this when we run the program.
- To view a single instruction run: 'x /i <address>'
 - Stands for eXamine instruction at <address>
- You can print multiple instructions via 'x /###i <address>' where ## is how many to print

gdb - Gnu Debugger

A look around

```
0x00000000000001181 <+4>:    push    rbp
0x00000000000001182 <+5>:    mov     rbp, rsp
0x00000000000001185 <+8>:    sub     rsp, 0x10
0x00000000000001189 <+12>:   mov     DWORD PTR [rbp-0x4], edi
0x0000000000000118c <+15>:   mov     QWORD PTR [rbp-0x10], rsi
0x00000000000001190 <+19>:   cmp     DWORD PTR [rbp-0x4], 0x1
0x00000000000001194 <+23>:   jle     0x11ab <main+46>
0x00000000000001196 <+25>:   mov     rax, QWORD PTR [rbp-0x10]
0x0000000000000119a <+29>:   add     rax, 0x8
0x0000000000000119e <+33>:   mov     rax, QWORD PTR [rax]
0x000000000000011a1 <+36>:   mov     rdi, rax
0x000000000000011a4 <+39>:   call    0x1149 <print_string>
0x000000000000011a9 <+44>:   jmp     0x11b7 <main+58>
0x000000000000011ab <+46>:   lea     rdi, [rip+0xe5e]          # 0x2010
0x000000000000011b2 <+53>:   call    0x1149 <print_string>
0x000000000000011b7 <+58>:   mov     eax, 0x0
0x000000000000011bc <+63>:   leave
0x000000000000011bd <+64>:   ret
End of assembler dump.
(gdb) x /i main
0x117d <main>:      endbr64
(gdb) x /i 0x117d
0x117d <main>:      endbr64
(gdb) █
```

gdb - Gnu Debugger

A look around

```
0x000000000000011a9 <+44>:    jmp     0x11b7 <main+58>
0x000000000000011ab <+46>:    lea     rdi,[rip+0xe5e]      # 0x2010
0x000000000000011b2 <+53>:    call   0x1149 <print_string>
0x000000000000011b7 <+58>:    mov     eax,0x0
0x000000000000011bc <+63>:    leave
0x000000000000011bd <+64>:    ret
End of assembler dump.
(gdb) x /i main
0x117d <main>:      endbr64
(gdb) x /i 0x117d
0x117d <main>:      endbr64
(gdb) x /5i main
0x117d <main>:      endbr64
0x1181 <main+4>:    push    rbp
0x1182 <main+5>:    mov     rbp, rsp
0x1185 <main+8>:    sub     rsp, 0x10
0x1189 <main+12>:   mov     DWORD PTR [rbp-0x4], edi
(gdb) x /5i 0x117d
0x117d <main>:      endbr64
0x1181 <main+4>:    push    rbp
0x1182 <main+5>:    mov     rbp, rsp
0x1185 <main+8>:    sub     rsp, 0x10
0x1189 <main+12>:   mov     DWORD PTR [rbp-0x4], edi
(gdb) █
```

gdb - Gnu Debugger

Breakpoints

- Breakpoints are ways to stop execution at different points in the program so that you may examine what is going on
- 'break *<address>|<symbol>'
- Let's set a breakpoint on the main function before we run it. This way we can see the execution of the function. If we just ran the application without a breakpoint it would just execute the entire thing
- Can be done two ways:
 - break *0x117d
 - break *main

gdb - Gnu Debugger

A look around

```
0x000000000000011b2 <+53>:    call    0x1149 <print_string>
0x000000000000011b7 <+58>:    mov     eax,0x0
0x000000000000011bc <+63>:    leave
0x000000000000011bd <+64>:    ret
End of assembler dump.
(gdb) x /i main
0x117d <main>:             endbr64
(gdb) x /i 0x117d
0x117d <main>:             endbr64
(gdb) x /5i main
0x117d <main>:             endbr64
0x1181 <main+4>:           push    rbp
0x1182 <main+5>:           mov     rbp, rsp
0x1185 <main+8>:           sub     rsp, 0x10
0x1189 <main+12>:          mov     DWORD PTR [rbp-0x4], edi
(gdb) x /5i 0x117d
0x117d <main>:             endbr64
0x1181 <main+4>:           push    rbp
0x1182 <main+5>:           mov     rbp, rsp
0x1185 <main+8>:           sub     rsp, 0x10
0x1189 <main+12>:          mov     DWORD PTR [rbp-0x4], edi
(gdb) break *main
Breakpoint 1 at 0x117d
(gdb) █
```

gdb - Gnu Debugger

A look around

- Now let's run the application
- type 'r'
- We want to be able to see the instructions as we step through so run
 - `display /i $rip`
 - This is a command that runs after each step to print the instruction that is pointed to by the RIP register.

gdb - Gnu Debugger

A look around

```
(gdb) x /i 0x117d
0x117d <main>:      endbr64
(gdb) x /5i main
0x117d <main>:      endbr64
0x1181 <main+4>:    push    rbp
0x1182 <main+5>:    mov     rbp, rsp
0x1185 <main+8>:    sub     rsp, 0x10
0x1189 <main+12>:   mov     DWORD PTR [rbp-0x4], edi
(gdb) x /5i 0x117d
0x117d <main>:      endbr64
0x1181 <main+4>:    push    rbp
0x1182 <main+5>:    mov     rbp, rsp
0x1185 <main+8>:    sub     rsp, 0x10
0x1189 <main+12>:   mov     DWORD PTR [rbp-0x4], edi
(gdb) break *main
Breakpoint 1 at 0x117d
(gdb) r
Starting program: /thayerfs/home/f00543d/recourse/hello_world

Breakpoint 1, 0x00005555555517d in main ()
(gdb) display/i $rip
1: x/i $rip
=> 0x5555555517d <main>:      endbr64
(gdb) █
```

gdb - Gnu Debugger

A look around

- Step through a few instructions by typing 'si'
- If you type it once then continue hitting enter then gdb will repeat your previous command.

```
(gdb) r
Starting program: /thayerfs/home/f00543d/recourse/hello_world

Breakpoint 1, 0x00005555555517d in main ()
(gdb) display/i $rip
1: x/i $rip
=> 0x5555555517d <main>:      endbr64
(gdb) si
0x000055555555181 in main ()
1: x/i $rip
=> 0x55555555181 <main+4>:   push    rbp
(gdb)
0x000055555555182 in main ()
1: x/i $rip
=> 0x55555555182 <main+5>:   mov     rbp, rsp
(gdb)
0x000055555555185 in main ()
1: x/i $rip
=> 0x55555555185 <main+8>:   sub     rsp, 0x10
(gdb)
0x000055555555189 in main ()
1: x/i $rip
=> 0x55555555189 <main+12>:  mov     DWORD PTR [rbp-0x4], edi
(gdb) █
```

gdb - Gnu Debugger

Short aside: ASLR

- Address Space Layout Randomization
- Many years ago the address space wasn't randomised. An application that was executed on one computer would have the same address space each time it ran on any computer it ran on.
- See any problems with this?
- Now: By default, programs are loaded at random addresses. Also, the stack and heap spaces are usually randomised
- Caveat, when you launch a program from gdb ASLR is often turned off.

gdb - Gnu Debugger

A look around

- Check out the register state: 'i r' short for 'info reg'
- What is rdi? rsi?

```
rax      0x5555555517d  93824992235901
rbx      0x0          0
rcx      0x555555551c0  93824992235968
rdx      0x7fffffff5c8  140737488348616
rsi      0x7fffffff5b8  140737488348600
rdi      0x1          1
rbp      0x7fffffff4d0  0x7fffffff4d0
rsp      0x7fffffff4c0  0x7fffffff4c0
r8       0x7ffff7dced80 140737351839104
r9       0x7ffff7dced80 140737351839104
r10      0x2          2
r11      0x1f         31
r12      0x55555555060  93824992235616
r13      0x7fffffff5b0  140737488348592
r14      0x0          0
r15      0x0          0
rip      0x55555555189  0x55555555189 <main+12>
eflags   0x206        [ PF IF ]
cs       0x33         51
ss       0x2b         43
ds       0x0          0
es       0x0          0
fs       0x0          0
---Type <return> to continue, or q <return> to quit---
```

gdb - Gnu Debugger

A look around

- Examine Memory:
- 'x /gx'
 - 'x' Says that we want to look at something
 - 'g' indicates the size. Could be 'g', 'w', 'h', 'b' for 64-,32-,16-, and 8-bits
 - 'x ' display as hex digits. You can also specify other types such as 's' for string

```
es          0x0      0
fs          0x0      0
---Type <return> to continue, or q <return> to quit---
gs          0x0      0
k0          0x0      0
k1          0x0      0
k2          0x0      0
k3          0x0      0
k4          0x0      0
k5          0x0      0
k6          0x0      0
k7          0x0      0
(gdb) x /gx $rsi
0x7fffffff5b8: 0x00007fffffff839
(gdb) x /2gx $rsi
0x7fffffff5b8: 0x00007fffffff839      0x0000000000000000
(gdb) x /gx 0x00007fffffff839
0x7fffffff839: 0x667265796168742f
(gdb) x /s 0x00007fffffff839
0x7fffffff839: "/thayerfs/home/f00543d/recourse/hello_world"
(gdb) x /16bx 0x00007fffffff839
0x7fffffff839: 0x2f      0x74      0x68      0x61      0x79      0x65      0x72      0x66
0x7fffffff841: 0x73      0x2f      0x68      0x6f      0x6d      0x65      0x2f      0x66
(gdb) █
```

gdb - Gnu Debugger

A look around

- Single step until the comparison
- Examine the memory referenced by that address
- What are we looking at?

```
(gdb) x /gx $rsi
0x7fffffff5b8: 0x00007fffffff839
(gdb) x /2gx $rsi
0x7fffffff5b8: 0x00007fffffff839      0x0000000000000000
(gdb) x /gx 0x00007fffffff839
0x7fffffff839: 0x667265796168742f
(gdb) x /s 0x00007fffffff839
0x7fffffff839: "/thayerfs/home/f00543d/recourse/hello_world"
(gdb) x /16bx 0x00007fffffff839
0x7fffffff839: 0x2f      0x74      0x68      0x61      0x79      0x65      0x72      0x66
0x7fffffff841: 0x73      0x2f      0x68      0x6f      0x6d      0x65      0x2f      0x66
(gdb) si
0x00005555555518c in main ()
1: x/i $pc
=> 0x5555555518c <main+15>:      mov     QWORD PTR [rbp-0x10],rsi
(gdb) si
0x000055555555190 in main ()
1: x/i $pc
=> 0x55555555190 <main+19>:      cmp     DWORD PTR [rbp-0x4],0x1
(gdb) x /dx $rbp-4
0x7fffffff4cc: 0x01
(gdb) x /wx $rbp-4
0x7fffffff4cc: 0x00000001
(gdb) █
```


gdb - Gnu Debugger

A look around

- Check out the flags that are set prior to the comparison.
- Single step then look at them again.
- Did anything change?
- If so what?
- Will this jump be taken?

```
(gdb) x /16bx 0x00007fffffff839
0x7fffffff839: 0x2f  0x74  0x68  0x61  0x79  0x65  0x72  0x66
0x7fffffff841: 0x73  0x2f  0x68  0x6f  0x6d  0x65  0x2f  0x66
(gdb) si
0x00005555555518c in main ()
1: x/i $pc
=> 0x5555555518c <main+15>:  mov    QWORD PTR [rbp-0x10],rsi
(gdb) si
0x000055555555190 in main ()
1: x/i $pc
=> 0x55555555190 <main+19>:  cmp    DWORD PTR [rbp-0x4],0x1
(gdb) x /dx $rbp-4
0x7fffffff84cc: 0x01
(gdb) x /wx $rbp-4
0x7fffffff84cc: 0x00000001
(gdb) i r $eflags
eflags          0x206    [ PF IF ]
(gdb) si
0x000055555555194 in main ()
1: x/i $pc
=> 0x55555555194 <main+23>:  jle    0x555555551ab <main+46>
(gdb) i r $eflags
eflags          0x246    [ PF ZF IF ]
(gdb) █
```

gdb - Gnu Debugger

A look around

- Single step and you see that the jump is taken.
- The next instruction show an interesting way to address a section in memory called “rip-relative addressing”
- It happens when you need to access something where the address is unknown at compile time but the location relative to the instruction is known.

```
0x00005555555518c in main ()
1: x/i $pc
=> 0x5555555518c <main+15>:    mov     QWORD PTR [rbp-0x10],rsi
(gdb) si
0x000055555555190 in main ()
1: x/i $pc
=> 0x55555555190 <main+19>:    cmp     DWORD PTR [rbp-0x4],0x1
(gdb) x /dx $rbp-4
0x7fffffffef4cc: 0x01
(gdb) x /wx $rbp-4
0x7fffffffef4cc: 0x00000001
(gdb) i r $eflags
eflags          0x206    [ PF IF ]
(gdb) si
0x000055555555194 in main ()
1: x/i $pc
=> 0x55555555194 <main+23>:    jle     0x555555551ab <main+46>
(gdb) i r $eflags
eflags          0x246    [ PF ZF IF ]
(gdb) si
0x0000555555551ab in main ()
1: x/i $pc
=> 0x555555551ab <main+46>:    lea     rdi,[rip+0xe5e]      # 0x555555556010
(gdb) █
```

gdb - Gnu Debugger

A look around

- Single step and let's see what rdi is pointing to.
- We are about to execute a call instruction. Can someone explain what this instruction will do?
- What will the return address be?

```
=> 0x55555555190 <main+19>:    cmp    DWORD PTR [rbp-0x4],0x1
(gdb) x /dx $rbp-4
0x7fffffffef4cc: 0x01
(gdb) x /wx $rbp-4
0x7fffffffef4cc: 0x00000001
(gdb) i r $eflags
eflags          0x206    [ PF IF ]
(gdb) si
0x000055555555194 in main ()
1: x/i $pc
=> 0x55555555194 <main+23>:    jle     0x555555551ab <main+46>
(gdb) i r $eflags
eflags          0x246    [ PF ZF IF ]
(gdb) si
0x0000555555551ab in main ()
1: x/i $pc
=> 0x555555551ab <main+46>:    lea     rdi,[rip+0xe5e]      # 0x555555556010
(gdb) si
0x0000555555551b2 in main ()
1: x/i $pc
=> 0x555555551b2 <main+53>:    call   0x55555555149 <print_string>
(gdb) x /s $rdi
0x555555556010: "Hello World!"
(gdb) █
```


gdb - Gnu Debugger

A look around

- 0x5555555551b7
- 'bt' backtrace shows the call stack
- Single step until you reach the call to printf.
- Look at the function arguments. What are they?

```
1: x/i $pc
=> 0x555555555194 <main+23>:    jle    0x5555555551ab <main+46>
(gdb) i r $eflags
eflags          0x246    [ PF ZF IF ]
(gdb) si
0x00005555555551ab in main ()
1: x/i $pc
=> 0x5555555551ab <main+46>:    lea    rdi,[rip+0xe5e]      # 0x5555555556010
(gdb) si
0x00005555555551b2 in main ()
1: x/i $pc
=> 0x5555555551b2 <main+53>:    call  0x555555555149 <print_string>
(gdb) x /s $rdi
0x5555555556010: "Hello World!"
(gdb) si
0x0000555555555149 in print_string ()
1: x/i $pc
=> 0x555555555149 <print_string>:    endbr64
(gdb) x /gx $rsp
0x7fffffffef4b8: 0x00005555555551b7
(gdb) bt
#0  0x0000555555555149 in print_string ()
#1  0x00005555555551b7 in main ()
(gdb) █
```

gdb - Gnu Debugger

A look around

```
0x000055555555160 in print_string ()
1: x/i $pc
=> 0x55555555160 <print_string+23>:   mov     rax,QWORD PTR [rbp-0x8]
(gdb)
0x000055555555164 in print_string ()
1: x/i $pc
=> 0x55555555164 <print_string+27>:   mov     rsi,rax
(gdb)
0x000055555555167 in print_string ()
1: x/i $pc
=> 0x55555555167 <print_string+30>:   lea     rdi,[rip+0xe96]          # 0x555555556004
(gdb)
0x00005555555516e in print_string ()
1: x/i $pc
=> 0x5555555516e <print_string+37>:   mov     eax,0x0
(gdb)
0x000055555555173 in print_string ()
1: x/i $pc
=> 0x55555555173 <print_string+42>:   call   0x55555555050 <printf@plt>
(gdb) x /s $rdi
0x555555556004: "String: %s\n"
(gdb) x /s $rsi
0x555555556010: "Hello World!"
(gdb) █
```


gdb - Gnu Debugger

A look around

- We don't want to go into printf so instead of using "single step" use "ni" for next instruction
- This command "skips" the function call.
- It actually just executes it and stops after it returns.

```
1: x/i $pc
=> 0x55555555164 <print_string+27>:  mov    rsi, rax
(gdb)
0x000055555555167 in print_string ()
1: x/i $pc
=> 0x55555555167 <print_string+30>:  lea    rdi, [rip+0xe96]          # 0x555555556004
(gdb)
0x00005555555516e in print_string ()
1: x/i $pc
=> 0x5555555516e <print_string+37>:  mov    eax, 0x0
(gdb)
0x000055555555173 in print_string ()
1: x/i $pc
=> 0x55555555173 <print_string+42>:  call   0x55555555050 <printf@plt>
(gdb) x /s $rdi
0x555555556004: "String: %s\n"
(gdb) x /s $rsi
0x555555556010: "Hello World!"
(gdb) ni
String: Hello World!
0x000055555555178 in print_string ()
1: x/i $pc
=> 0x55555555178 <print_string+47>:  jmp    0x5555555517b <print_string+50>
(gdb) █
```

gdb - Gnu Debugger

A look around

- We made it through a single concrete execution
- Let's add a command line argument.
- Type 'r "This is a single argument"'
- Enter 'y' when it asks if you are sure.

```
(gdb)
0x0000555555551bc in main ()
1: x/i $pc
=> 0x555555551bc <main+63>:    leave
(gdb)
0x0000555555551bd in main ()
1: x/i $pc
=> 0x555555551bd <main+64>:    ret
(gdb)
__libc_start_main (main=0x5555555517d <main>, argc=1, argv=0x7fffffffe5b8, init=<optimized out>,
    fini=<optimized out>, rtld_fini=<optimized out>, stack_end=0x7fffffffe5a8)
    at ../csu/libc-start.c:344
344      ../csu/libc-start.c: No such file or directory.
1: x/i $pc
=> 0x7ffff7a03bf7 <__libc_start_main+231>:    mov     edi,eax
(gdb) r "This is a single argument"
The program being debugged has been started already.
Start it from the beginning? (y or n) y
Starting program: /thayerfs/home/f00543d/recourse/hello_world "This is a single argument"

Breakpoint 1, 0x00005555555517d in main ()
1: x/i $pc
=> 0x5555555517d <main>:      endbr64
(gdb) █
```

gdb - Gnu Debugger

A look around

- Take some time and do the same thing that we did before.
- Look at the flags before and after the CMP.
- Look at the return value of print_string
- Check out the arguments of printf()

gdb - Gnu Debugger

Memory Regions

- Make sure that you are still in gdb with a running application
- Run 'info proc mappings'
- Look for [stack] and [heap]

```
process 37716
Mapped address spaces:

   Start Addr           End Addr       Size     Offset objfile
   0x555555554000       0x555555555000     0x1000        0x0  /thayerfs/home/f00543d/recourse/hello_wor
ld
   0x555555555000       0x555555556000     0x1000     0x1000  /thayerfs/home/f00543d/recourse/hello_wor
ld
   0x555555556000       0x555555557000     0x1000     0x2000  /thayerfs/home/f00543d/recourse/hello_wor
ld
   0x555555557000       0x555555558000     0x1000     0x2000  /thayerfs/home/f00543d/recourse/hello_wor
ld
   0x555555558000       0x555555559000     0x1000     0x3000  /thayerfs/home/f00543d/recourse/hello_wor
ld
   0x7ffff79e2000       0x7ffff7bc9000    0x1e7000        0x0  /lib/x86_64-linux-gnu/libc-2.27.so
   0x7ffff7bc9000       0x7ffff7dc9000    0x200000    0x1e7000  /lib/x86_64-linux-gnu/libc-2.27.so
   0x7ffff7dc9000       0x7ffff7dcd000     0x4000    0x1e7000  /lib/x86_64-linux-gnu/libc-2.27.so
   0x7ffff7dcd000       0x7ffff7dcf000     0x2000    0x1eb000  /lib/x86_64-linux-gnu/libc-2.27.so
   0x7ffff7dcf000       0x7ffff7dd3000     0x4000        0x0
   0x7ffff7dd3000       0x7ffff7dfc000    0x29000        0x0  /lib/x86_64-linux-gnu/ld-2.27.so
   0x7ffff7fa6000       0x7ffff7fa8000     0x2000        0x0
   0x7ffff7ff5000       0x7ffff7ff7000     0x2000        0x0
   0x7ffff7ff7000       0x7ffff7ffa000     0x3000        0x0  [vvar]
---Type <return> to continue, or q <return> to quit---
```

gdb - Gnu Debugger

Memory Regions

- Make sure that you are still in gdb with a running application
- Run 'info proc mappings'
- Look for [stack].
- The pointer in rsp should be within this range.

```
process 37716
Mapped address spaces:

   Start Addr           End Addr       Size     Offset objfile
   0x555555554000       0x555555555000     0x1000        0x0  /thayerfs/home/f00543d/recourse/hello_wor
ld
   0x555555555000       0x555555556000     0x1000     0x1000  /thayerfs/home/f00543d/recourse/hello_wor
ld
   0x555555556000       0x555555557000     0x1000     0x2000  /thayerfs/home/f00543d/recourse/hello_wor
ld
   0x555555557000       0x555555558000     0x1000     0x2000  /thayerfs/home/f00543d/recourse/hello_wor
ld
   0x555555558000       0x555555559000     0x1000     0x3000  /thayerfs/home/f00543d/recourse/hello_wor
ld
   0x7ffff79e2000       0x7ffff7bc9000    0x1e7000        0x0  /lib/x86_64-linux-gnu/libc-2.27.so
   0x7ffff7bc9000       0x7ffff7dc9000    0x200000    0x1e7000  /lib/x86_64-linux-gnu/libc-2.27.so
   0x7ffff7dc9000       0x7ffff7dcd000     0x4000    0x1e7000  /lib/x86_64-linux-gnu/libc-2.27.so
   0x7ffff7dcd000       0x7ffff7dcf000     0x2000    0x1eb000  /lib/x86_64-linux-gnu/libc-2.27.so
   0x7ffff7dcf000       0x7ffff7dd3000     0x4000        0x0
   0x7ffff7dd3000       0x7ffff7dfc000    0x29000        0x0  /lib/x86_64-linux-gnu/ld-2.27.so
   0x7ffff7fa6000       0x7ffff7fa8000     0x2000        0x0
   0x7ffff7ff5000       0x7ffff7ff7000     0x2000        0x0
   0x7ffff7ff7000       0x7ffff7ffa000     0x3000        0x0  [vvar]
---Type <return> to continue, or q <return> to quit---
```

/proc

Memory Regions

- /proc is a file system that contains process information
- Open another terminal on the same babylon server.
- Run 'ps -aux | grep <netid> | grep hello_world' to get the process id

```
f00543d@babylon1:~$ ps aux | grep f00543d | grep hello_world
f00543d 32136 0.0 0.0 14436 1048 pts/12 S+ 16:02 0:00 grep hello_world
f00543d 37716 0.0 0.0 4400 888 pts/10 t 15:50 0:00 /thayerfs/home/f00543d/recourse/hello_
world This is a single argument
f00543d 43260 0.0 0.0 100324 43800 pts/10 S+ 15:04 0:00 gdb ./hello_world
f00543d@babylon1:~$
```


/proc

Memory Regions

- Run 'cat /proc/<pid>/maps | less'
- You get mostly the same information but the interesting bits are the permissions.
- r: read
- w: write
- x: execute
- p: private (copy on write)
- <https://www.kernel.org/doc/html/latest/filesystems/proc.html>

```
55555554000-55555555000 r--p 00000000 00:38 23537167833 /thayerfs/home/f00543d/recours
e/hello_world
55555555000-55555556000 r-xp 00001000 00:38 23537167833 /thayerfs/home/f00543d/recours
e/hello_world
55555556000-55555557000 r--p 00002000 00:38 23537167833 /thayerfs/home/f00543d/recours
e/hello_world
55555557000-55555558000 r--p 00002000 00:38 23537167833 /thayerfs/home/f00543d/recours
e/hello_world
55555558000-55555559000 rw-p 00003000 00:38 23537167833 /thayerfs/home/f00543d/recours
e/hello_world
7ffff79e2000-7ffff7bc9000 r-xp 00000000 09:7e 1311323 /lib/x86_64-linux-gnu/libc-2.2
7.so
7ffff7bc9000-7ffff7dc9000 ---p 001e7000 09:7e 1311323 /lib/x86_64-linux-gnu/libc-2.2
7.so
7ffff7dc9000-7ffff7dcd000 r--p 001e7000 09:7e 1311323 /lib/x86_64-linux-gnu/libc-2.2
7.so
7ffff7dcd000-7ffff7dcf000 rw-p 001eb000 09:7e 1311323 /lib/x86_64-linux-gnu/libc-2.2
7.so
7ffff7dcf000-7ffff7dd3000 rw-p 00000000 00:00 0
7ffff7dd3000-7ffff7dfc000 r-xp 00000000 09:7e 1310773 /lib/x86_64-linux-gnu/ld-2.27.
so
7ffff7fa6000-7ffff7fa8000 rw-p 00000000 00:00 0
7ffff7ff5000-7ffff7ff7000 rw-p 00000000 00:00 0
:
```

gdb - Gnu Debugger

On your own

- At this point you have the basic knowledge of how to dynamically reverse engineer an application
- There are more things to learn that are helpful
 - Conditional breakpoints
 - Modifying memory/registers
 - Scripting
 - I like using “gef” to make gdb a bit more useful

gdb - Gnu Debugger

Lab

- Open week2_day1_lab in Ghidra
- You can also use all the other tools that we have discussed
- What inputs are required to have the “SUCCESS!” string printed?

Homework

- Use the tools you know to figure out valid inputs
- This is an example of what we call CrackMes.
- They are designed to “legally” break software serial algorithms.
- For example, Windows required a serial number to verify that you had purchased the software. However, there existed an algorithm to verify what you entered. What if you just reverse engineered the algorithm and wrote some code that automatically generated serial numbers? Don't do this.
- CrackMes allow you to scratch the itch of breaking something but in a way where the feds won't show up at your door. :)
- There also exist legal ways to hack into servers that we will discuss at the end of the course.

Day 2 Homework

- You will be compiling and looking at the disassembly of 3 C programs
- If your machine isn't x86 you can log into the babylon servers and use those for compiling